

NATIONAL ENGINEERING CENTER

University of the Philippines
Diliman, Quezon City



7.0 Transformation and Loading Methodologies

Eugene Rex L. Jalao, Ph.D.

Associate Professor

Department Industrial Engineering and Operations Research

University of the Philippines Diliman

*Module 2 of the Business Intelligence and Analytics Track of
UP NEC and the UP Center of Business Intelligence*

Outline for This Training

1. Introduction to Data Warehousing
2. DW Lifecycle and Project Management
 - Case Study on DW PM
3. Dimensional Modeling
4. Designing Fact Tables
5. Designing Dimension Tables
 - Case Study on Dimension Modeling
6. Extraction Transformation and Loading
 - Case Study on ETL Planning
- 7. Transformation and Loading Methodologies**
 - Case Study on ETL**



Outline for This Session

- Transformation Process
- Dimension Table Processing
- Dimension Update Types
- Fact Table Processing
- Loading Process
- Case Study



Data Transformation

- Extracted data is **raw data** and it cannot be applied to the data warehouse
- All the extracted data must be made **usable** in the data warehouse.

Quality of data

- Major effort within data transformation is the improvement of **data quality**.
- This includes filling in the **missing values** for attributes in the extracted data.
- Data quality is of paramount importance in the data warehouse because the effect of strategic decisions based on **incorrect information** can be devastating.

Basic Tasks in Data Transformation

- Selection
 - Select either **whole** records or **parts** of several records from the source systems.
- Splitting/joining –
 - Sometimes (uncommonly), you will be **splitting** the selected parts even further during data transformation.
 - **Joining** of parts selected from many source systems is more widespread in the data warehouse environment.



Basic Tasks in Data Transformation

- Conversion
 - It includes a large variety of rudimentary conversions of single fields for two primary reasons
 - one to **standardize** among the data extractions from disparate source systems,
 - To make the fields **usable and understandable** to the users.
- Summarization.
 - Sometimes it is **not feasible** to keep data at the lowest level of detail in the data warehouse.
 - It may be that **none of the users** ever need data at the lowest granularity for analysis or querying.



Basic Tasks in Data Transformation

- Enrichment
 - Rearrangement and simplification of individual fields to **make them more useful** for the data warehouse environment.
 - You may use one or more fields from the same input record to create **a better view** of the data for the data warehouse.
 - This principle is extended when one or more fields **originate from multiple records**, resulting in a single field for the data warehouse.



Major Transformation Types

- Format Revisions
- Decoding of Fields
- Calculated and Derived Values
- Splitting of Single Fields
- Merging of Information
- Character Set Conversion
- Conversion of Units of Measurements
- Date/Time Conversion
- Summarization
- Key Restructuring
- Deduplication



Outline for This Session

- Transformation Process
- **Dimension Table Processing**
- Dimension Update Types
- Fact Table Processing
- Loading Process
- Case Study



ETL Intermediate Dimension Tables

- Extract Tables (S, X)
- Master Tables (M)
- Change and Error Tables (C, E)
- Transform Tables (T)
- Insert Tables (I)
- Update Tables (U)
- Dimension Tables (D, H)

Dimension and Intermediate Table Naming Conventions

- S_Customers
- M_Customers
- X_Customers
- C_Customers
- E_Customers
- T_Customer
- I_Customer
- U_Customer
- D_Customer
- H_Customer_Audit_History



Full Extract Table (S)

- Allowable and non-allowable **nulls** accepted
- Source Logical Primary Key **disabled**
- Unique description and name constraints **disabled**
- Contains data directly extracted from **source system**
- Matched against **master to detect changes**
- Non-enforced constraints are enabled in **change table**



Extract Master Table (M)

- Used to detect changes if it is a **complete** rather than **incremental** extract
- Used to verify **additions** if it is an incremental extract

Incremental Extract Table (X)

- Contains data directly extracted from source system
- Allowable nulls changed to 'Missing Optional'
- Logical Primary Key enabled
- Unique description and name constraints enabled
- Non-allowable nulls either rejected or flagged
- As many constraints as possible should be applied here
 - Entity integrity – logical primary key
 - Missing required data – not null or blanks
- This is the place to catch as many errors as early as possible
 - Makes the rest of the processing cleaner
 - Minimizes work needed to be done



Change Table (C)

- Used to hold **new and changed records** from the match process in the case of full extracts
- Not used for incremental extracts

Transform Table (T)

- Holds **new and changed dimension rows** after transforms have been applied
- Has the **structure of the dimension table** where this data will eventually reside
- Up to this point all the intermediate tables had the structure of the source system files/tables

Transform Tables

- Insert Table (I)
 - Holds **new dimension rows** before being added to dimension
- Update table (U)
 - Holds **changed dimension rows** before being added to dimension

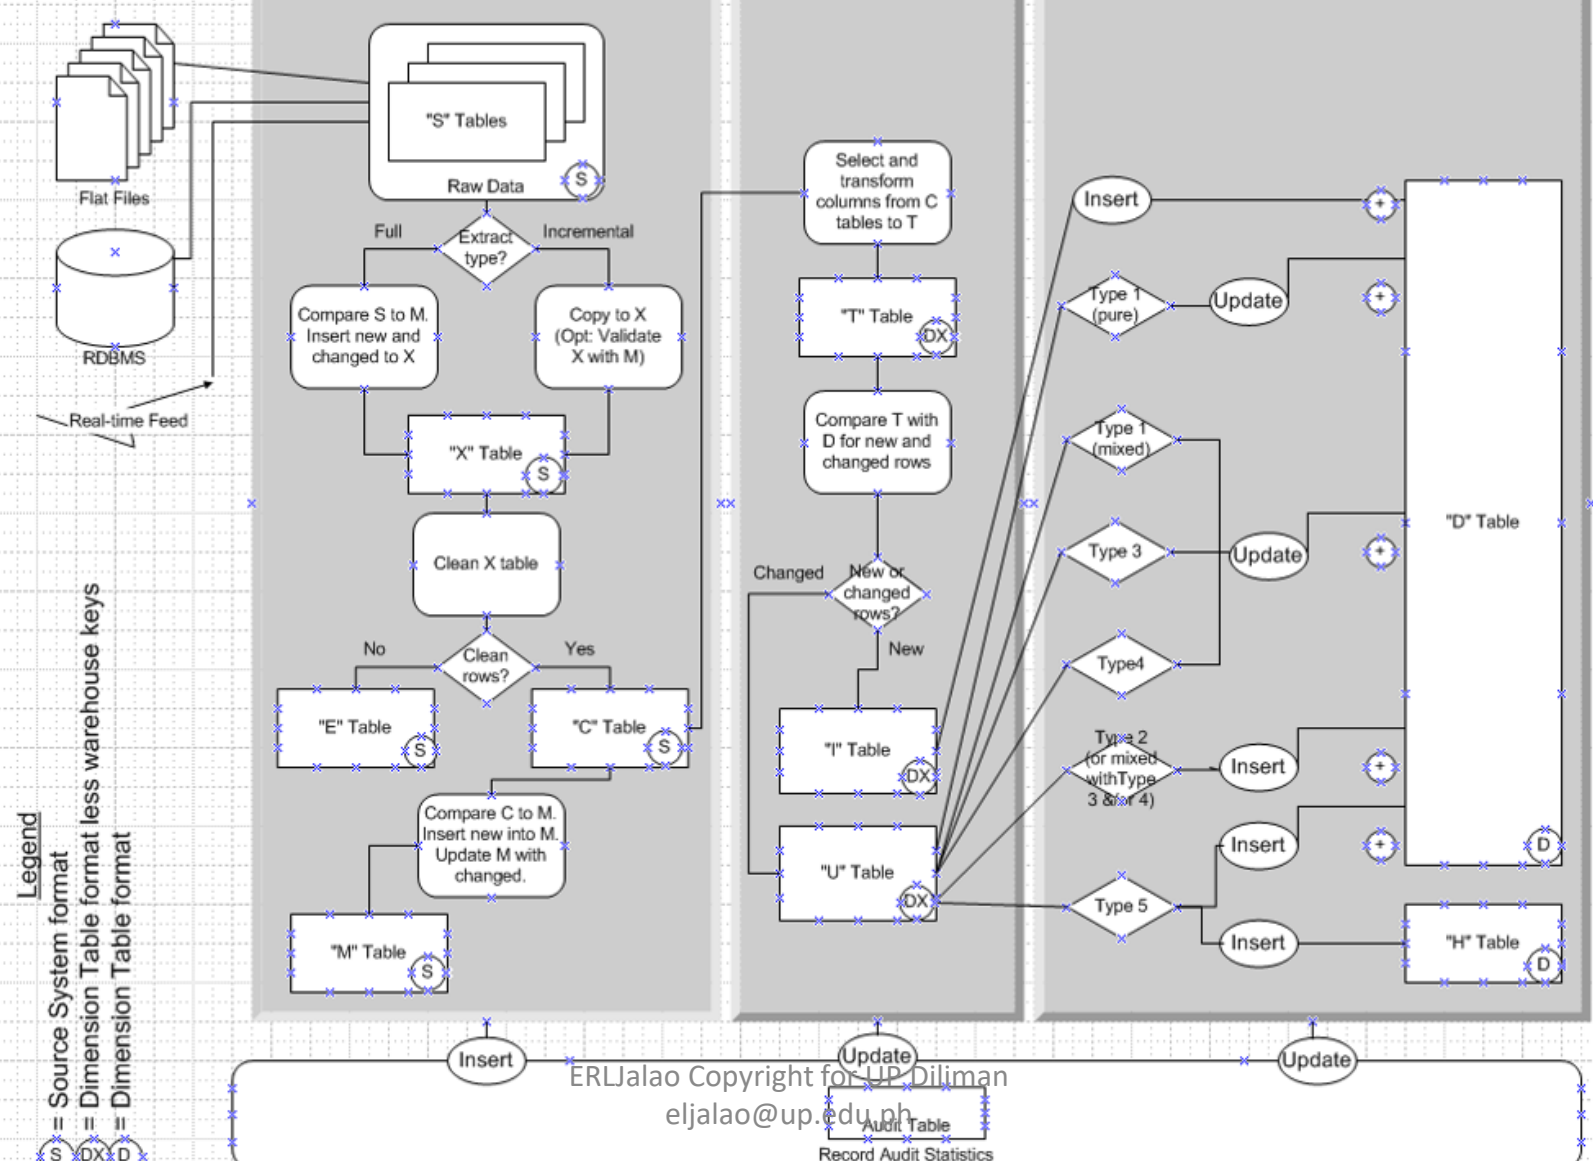
Dimension Audit History Table (H)

- Holds **complete history** of all detail and analytic attributes of a dimension
- Used if the dimension needs to be analyzed for **auditing purposes**
- Also known as a **dimension history table**
- Should only be used if need to use the dimension history table as a fact table
- Not joined to a fact table

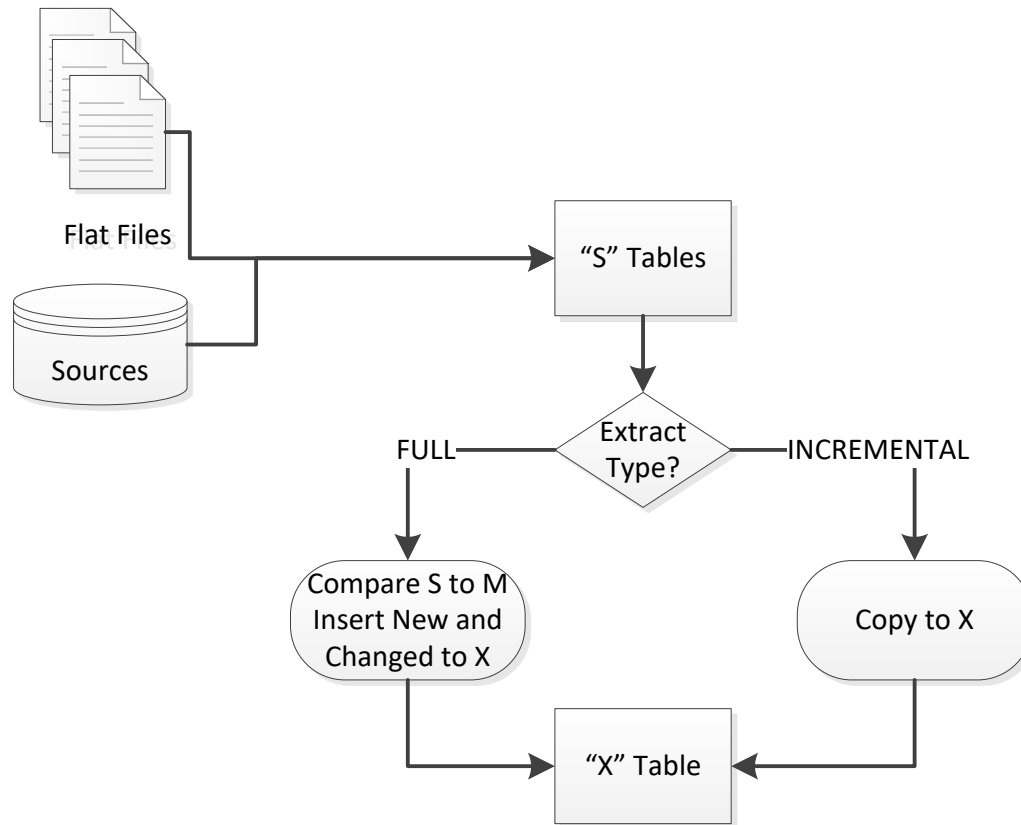
Process Overview

ETL ROADMAP

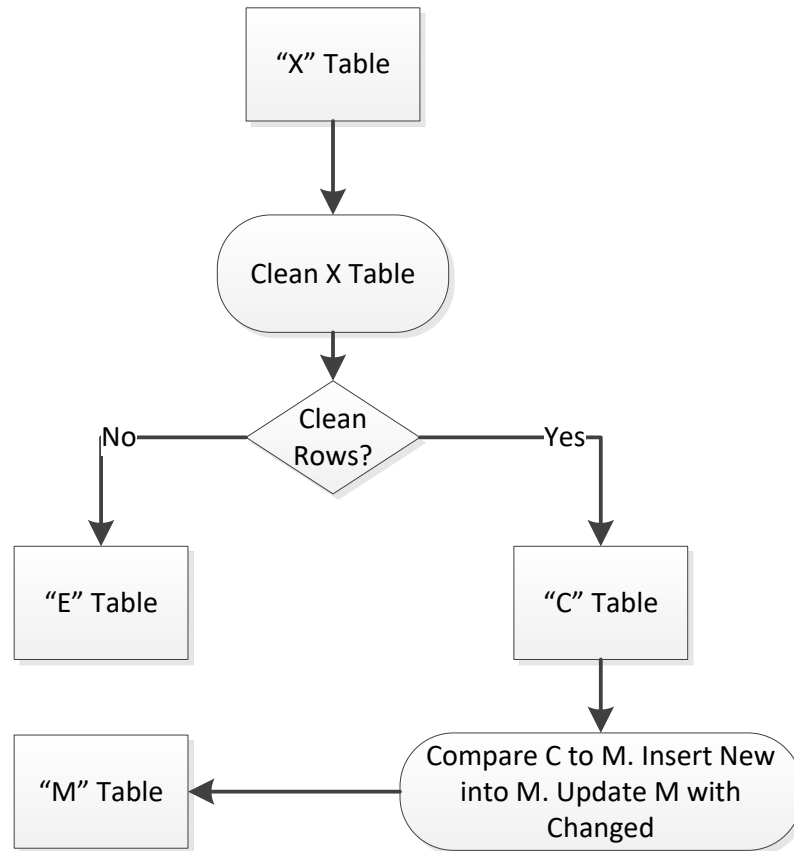
Version 5



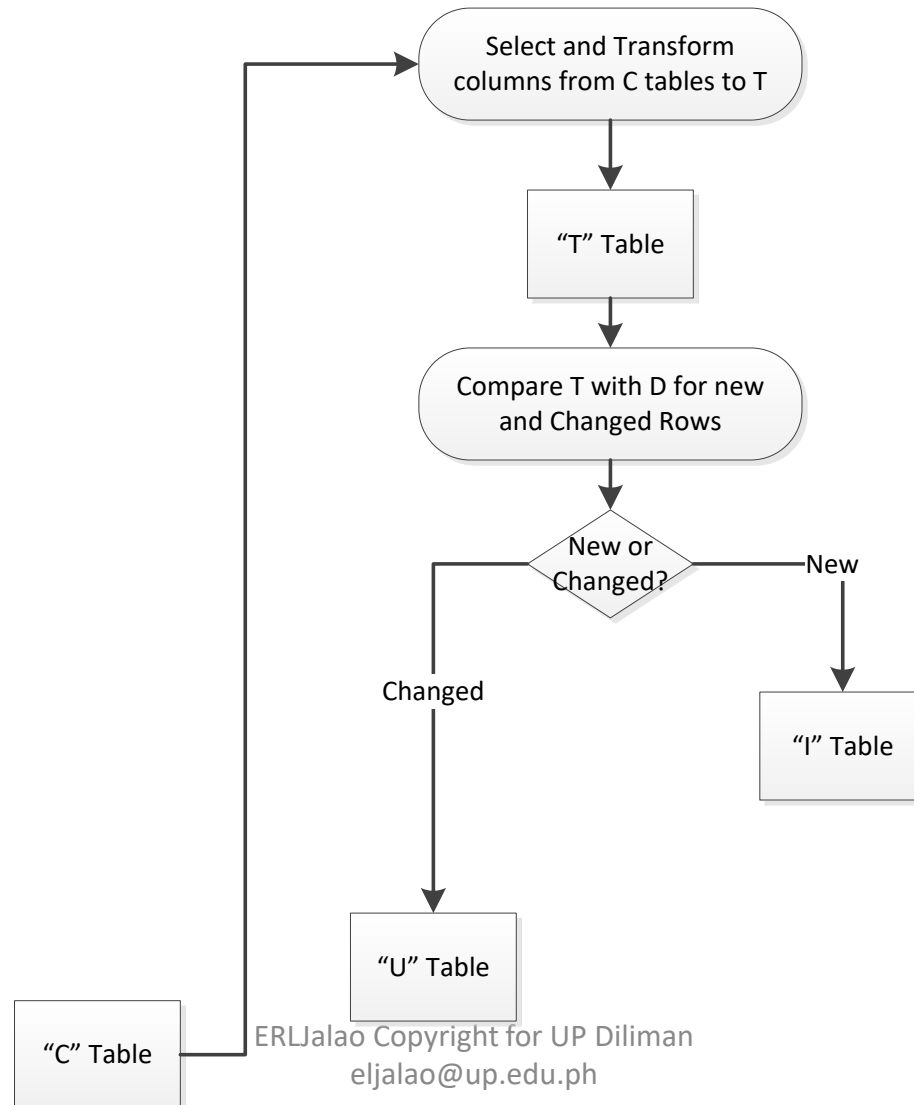
Step 1 – Source to Extract Table



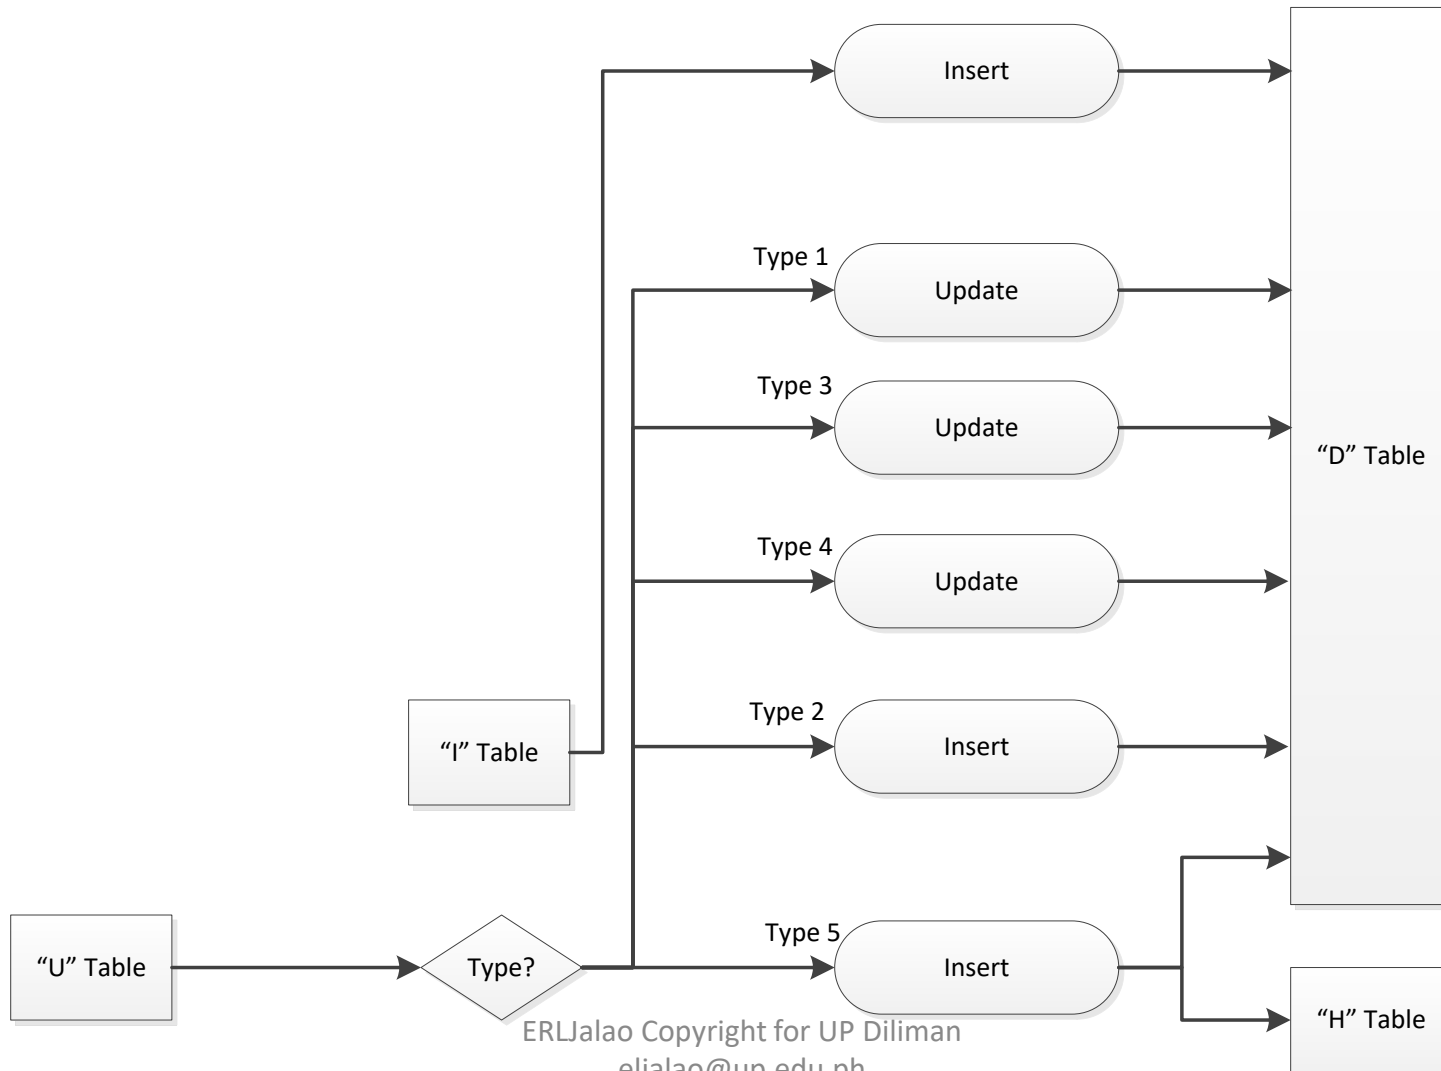
Step 2 – Find or Verify New & Changed



Step 3 – Insert into Respective Intermediate Table



Step 4 and 5



Full Extract Processing up to this point

- Load **Source Table**
- Minus (take set difference between) **Source Table** and **Master Table** to detect new and changed
- Insert clean rows into **Change table**, dirty rows into error table
- Transform columns from **Change Tables** to **Transform tables**
- Outer Join to dimension table to separate new and changed
- Insert new into **Insert Table**
- Insert changed into **Update Table**
- Insert into **Data Warehouse**



Incremental Extract Processing up to this point

- Load **Extract Table**
- Minus extract and master to 'ensure' only new and changed rows
- Insert clean rows into **Change Table**, dirty rows into error table
- **Transform Columns** from **Change Tables** to **Transform Tables**
- Outer Join to dimension table to separate new and changed
- Insert new into **Insert Table**
- Insert changed into **Update Table**



Outline for This Session

- Transformation Process
- Dimension Table Processing
- **Dimension Update Types**
- Fact Table Processing
- Loading Process
- Case Study



New Dimension Row Processing

- Apply transformations
- Generate new key
- Bulk load into dimension
- Insert into Dimension Audit History



Dimension Update Processing

- Apply transformations
- If a dimension table contains all type 1 columns then
- Simply – apply the updates
- In order
 - Use type 1 changed values
 - Handle type 3 changes
 - Handle type 4 changes
 - Handle type 2 changes



Dimension Attribute History

- Introduction to Dimension Attribute History
- Dimension Attribute History Type 1
- Dimension Attribute History Type 2
- Dimension Attribute History Type 4
- Dimension Attribute History Type 3
- Dimension Attribute History Type 5

Tracking Dimension Table Attribute History

- Five Considerations
- Five Methods
 - Ignore History
 - Overwrite
 - Track History
 - New Row
 - New Column
 - New Dimension
 - Dimension History Fact Table
- Many Combinations



Six Basic Considerations

- The reason the change occurred, is it a **correction** to a wrong value?
- Dimension table size
- The **number of changes** that need tracking
- The **frequency** of the changes
- Is the attribute **analytical** (not unique) or **detail** (unique)?
- Is the requirement **analytical** or **audit** based?



Dimension Attributes

- Change
 - Slowly
 - Quickly
 - In mass
- Lack of history can cause
 - Incorrect analyses
 - Erroneous decisions
- Tracking history should **not** affect
 - Analytical performance
 - Ease of use



Formal Methodology

- Based on Business Requirements
- Documented
- Signed off by Business Users
- Administered by Data Authority/Steward



Assess Requirements and Change Profile for Each Dimension Attribute

Customer Location	Analytical or Detail	Change Profile	History Requirement
Customer_Locn_Id	Primary key	N/A	N/A
Business Name	Detail	Seldom	None
Type of Business	Analytical	Seldom	All
Number of Lines	Analytical	Frequent	All
Marketing Segment	Analytical	Moderate	All
Sales Org	Analytical	Moderate	All
Sales Person	Analytical	Frequent	All
Sales Region	Analytical	Moderate	Limited (current, previous)
Street Address	Detail	Moderate	None
Address Line 2	Detail	Moderate	None
Address Line 3	Detail	Moderate	None
City	Analytical	Moderate	All
State	Analytical	Moderate	All
Zip Code	Analytical	Moderate	All
Product	Analytical or Detail	Change Profile	History Requirement
Product_Id	Primary key	Seldom	
Product_Desc	Detail	Seldom	None
Product_Line	Analytical	Seldom	Limited (current, previous)
Design_Class	Analytical	Seldom	All

Tracking Methods

- Ignore History
 - Type 1 - overwrite value
- Track History
 - Type 2 - create new row with updated values
 - Type 3 - use columns to store limited history
 - Type 4 - use analytical dimensions
 - Type 5 - use dimension history fact table



No History Required

- Type 1 - **Overwrite** old value with new value
- Always used for **corrections**
- When to use it?
 - If attribute history is not of **analytical** or **historical** importance
 - If the update is a **correction**

Type 1 - Example

- UPDATE customer_dim SET
street_address TO '2323 Jennings
St.'
- UPDATE customer_dim SET
phone_number TO '258-1913'

All Changes Need To Be Tracked

- All changes are **tracked**
- Conditions:
 - The dimension is **not too large**
 - The changes are **not too frequent**
- Perfectly partitions **history**
- If one or more attributes are classified as Type 2, we classify the dimension as a **Type 2** dimension.
- Does not mean that all attributes must be **Type 2**
- Add a **new row** to the dimension each time one or more tracked attributes change



Type 2 - Create Another Dimension Row with New Warehouse Key

Sales Date Dim
sales date key
sales date
sales month
sales calendar year

Sales Item Fact
sales date key
product key
qty
unit price
extended amount

Sales Fact sample data				
sales date key	product key	quantity	unit price	extended amount
23	10	25	5	125
73	1028	15	4	60

Product Dim
product key
product id
product desc
product weight
product size
version design class
propagated design class
effective date
current row indicator

```
SELECT product_name, sum(extended_amount)
FROM sales fact sf, sales_date_dim sd,
product_dim pd
WHERE join-conditions
AND product_id = '1256238'
```

Result:

product_name	sum
turbo xxx	185

Product Dim Sample Data

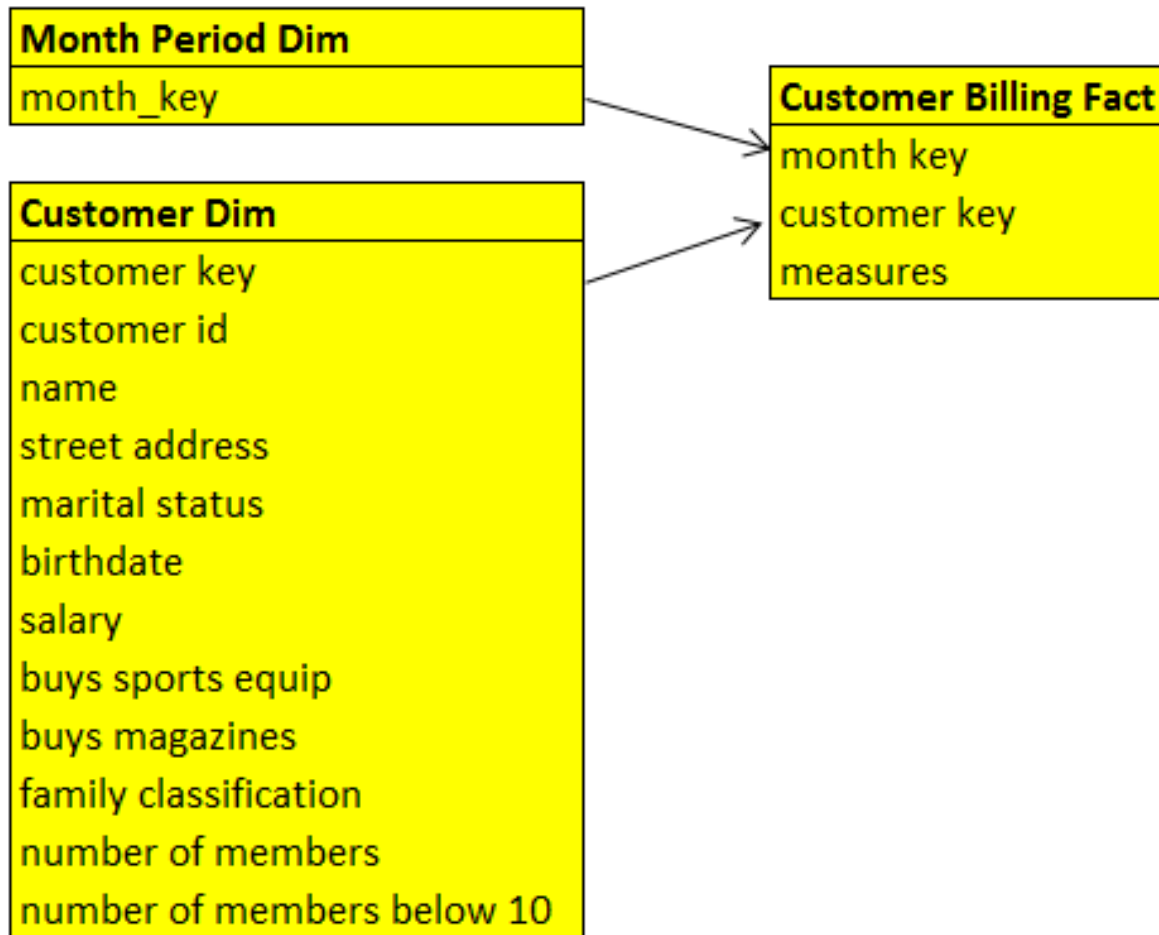
product key	product id	product desc	product weight	product size	version design class	propagated design class	effective date	current row indicator
10	1256238	turbo xxx	20	6	A	M	12/1/1990	N
1028	1256238	turbo xxx	30	6	M	M	4/10/2000	Y

All Changes Need to be Tracked

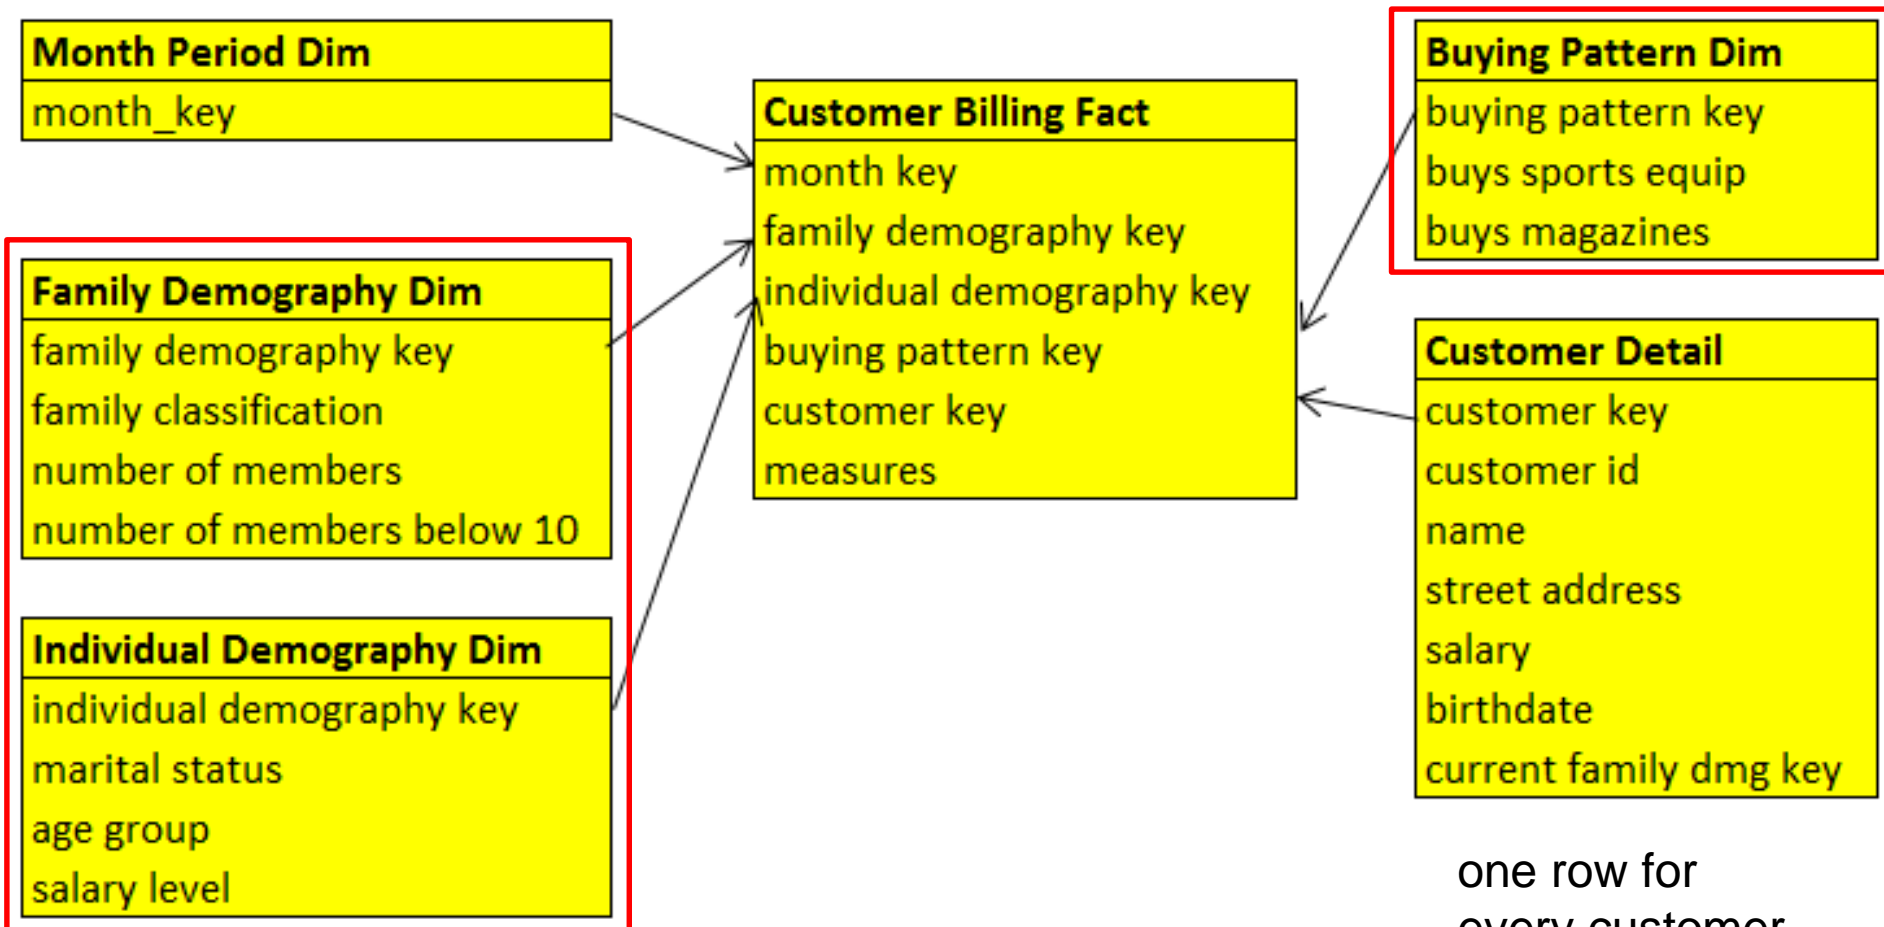
- Frequent Changes
- **Large table** (upwards of 100,000 rows or more, rather than a smaller table)
- Use **Type 4: Analytical Dimensions**
- Other Names:
 - Profile dimensions
 - Mini-dimensions
 - Sub-dimensions

Type 4 Example

Customer Dimension



Type 4 - Create Analytical Dimensions



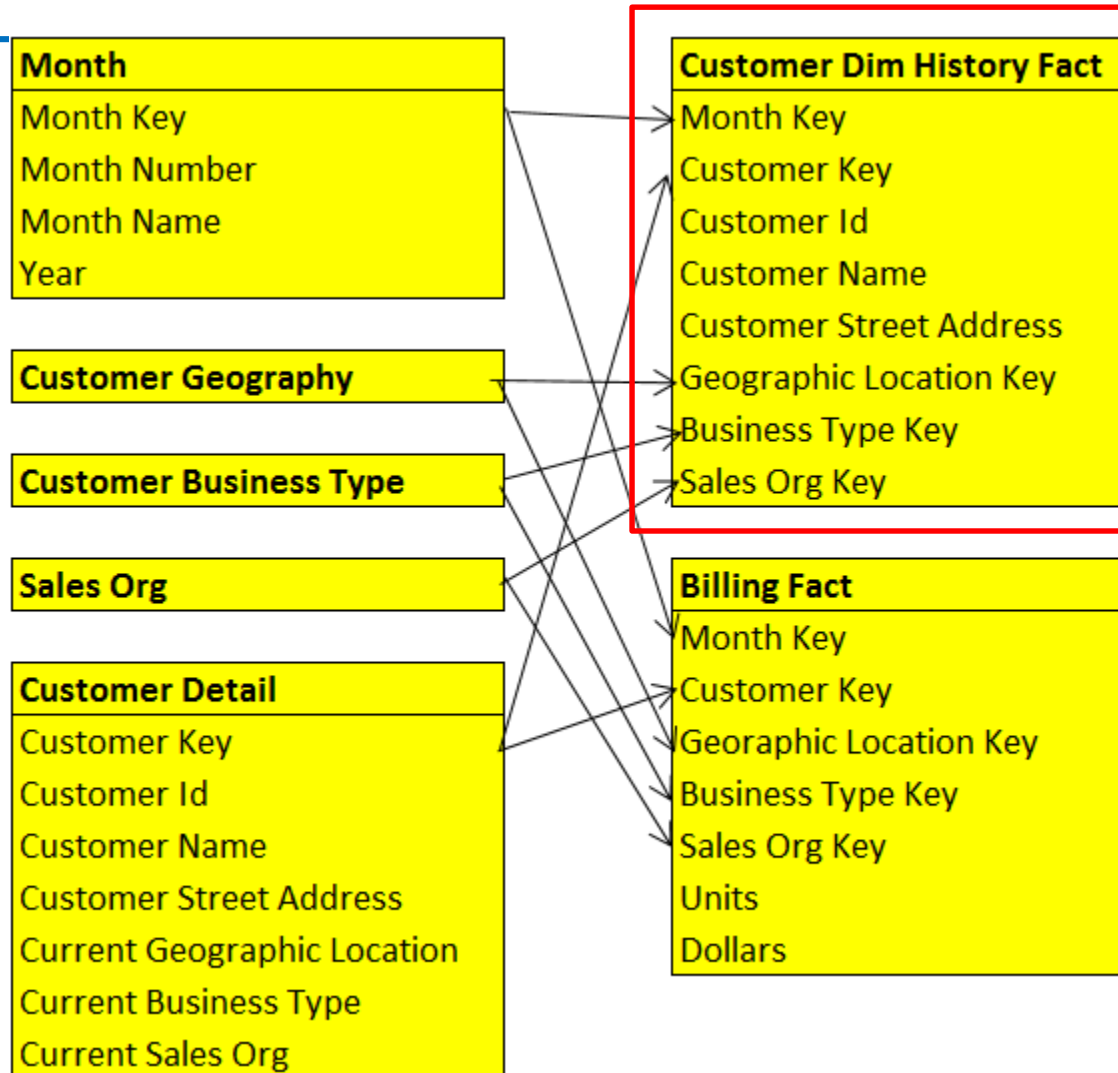
Automatic Fact Table History

- The fact table has **keys** to the customer attributes which were in effect when the fact table row was produced
- **Dimension history fact table** can be used to analyze dimension only history of all dimension attribute changes

Detail Attribute History and Dimension Change Tracking

- Need street address history
- Use Type 5 – Dimension History Fact Table
- It also supplements Type 4 Analytical Dimensions

Type 5 Dimension History Fact Table



Type 5 Summary


- Traps detail attribute **history**
- Should only be used to analyze **dimension-only history**
- Shouldn't be used **alone**
- Should not be used to query transaction or event history - use other types for that

Only a Limited Number of Changes Need Tracked


- Current versus Last
- Current and Planned
- Type 3 Examples
 - Analysis on new alignment versus old alignment
 - Analysis on planned realignment
- Only a limited number of changes need tracked
 - Not elegant
 - But very effective and high performing

Example

- Original Tables



- New Sales Person table with added column




Sample Queries

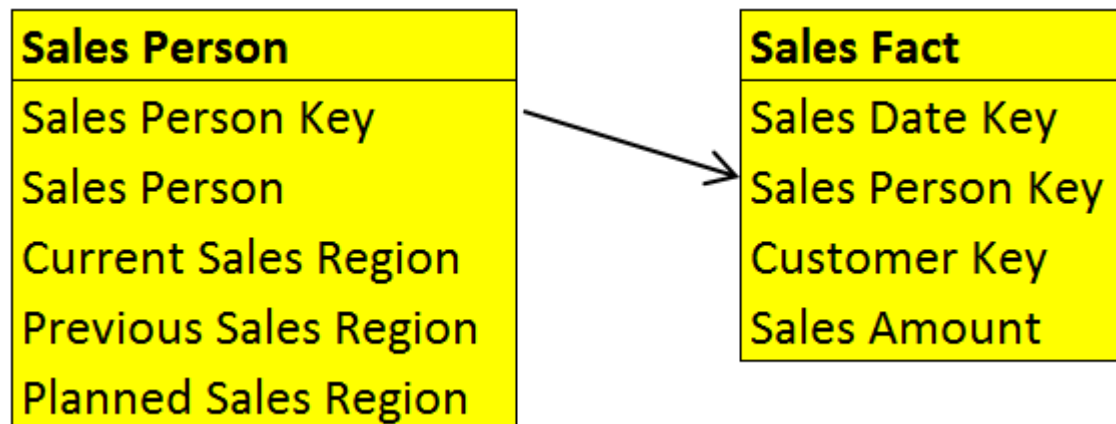
- Show me last month sales totals by new regions
 - `SELECT current_sales_region,
SUM(sales_amount) FROM sales_fact
WHERE month = 'LAST'`
- Show me last month sales totals by the old regions
 - `SELECT previous_sales_region,
SUM(sales_amount) FROM sales_fact
WHERE month = 'LAST'`

Example

- Original Tables



- New Sales Person table with added column



Sample Queries

- Show me last month sales totals by the proposed regions
 - `SELECT planned_sales_region,
SUM(sales_amount) FROM sales_fact
WHERE month = 'LAST'`

Choosing Dimension Attribute History Types

- Decision Matrix

History Requirement	History Strategy
None	Type 1
Limited	Type 3
All	Type 2,4,5

All	Analytical	small size - seldom, moderate or frequent chg	Type2
All		medium size - seldom and moderate chg	Type 2
All		large size - seldom and doesn't fit mini	Type 2
All		medium size - frequent chg	Type 4
All		large size - moderate and frequent chg	Type 4
All	Detail	all sizes	Type 5

Special Considerations

- Restating history
 - Essentially this means overriding history and replacing old values.
- Showing original or as of values
 - show both the original or 'as of values' and the new values

You Want History Restated

- Standard Type 2 attribute (small to moderate table, infrequent attribute changes)
- A product is now class B and you want the history restated
- Use Type 1 update instead of the standard Type 2

You Want History Restated ...

- Standard Type 4 attribute (large dimension or frequent changes)
- Customer Sales Person
 - Sales Person Dim
 - Customer Detail Dim
 - Key must be updated in the fact table

As of Values Required

- Who was the original sales person for the customer
 - Put attribute in Customer Detail Dimension
 - Original Sales Person Name

Example of Type Selection

Column	Type - Analytical/Detail	Frequency of Change	History Requirement
product_key	warehouse key	NA	NA
product_id	legacy key	doesn't change	NA
product_desc	detail business attribute	seldom	No history required
product_line_key	hierarchial warehouse key	NA	NA
product_line_code	analytical business attribute	moderate	only current and last
product_line_desc	analytical business attribute	moderate	only current and last
product_group_key	hierarchial warehouse key	NA	NA
product_group_code	analytical business attribute	moderate	only current and last
product_group_desc	analytical business attribute	moderate	only current and last
manufacturer_key	hierarchial warehouse key	NA	NA
manufacturer_code	analytical business attribute	seldom	all history
manufacturer_desc	analytical business attribute	seldom	all history
design_class_key	hierarchial warehouse key	NA	NA
design_class_code	analytical business attribute	seldom	all history
design_class_desc	analytical business attribute	seldom	all history



Decide on Method

Column	Type - Analytical/Detail	Change Frequency	History Requirement	History Strategy
product_key	warehouse key	NA	NA	NA
product_id	legacy key	doesn't change	NA	NA
product_desc	detail business attribute	seldom	No history required	Type 1
product_line_key	hierarchial warehouse key	NA	NA	NA
product_line_code	analytical business attribute	moderate	only current and last	Type 3 current & last
product_line_desc	analytical business attribute	moderate	only current and last	Type 3 current & last
product_group_key	hierarchial warehouse key	NA	NA	NA
product_group_code	analytical business attribute	moderate	only current and last	Type 3 current & last
product_group_desc	analytical business attribute	moderate	only current and last	Type 3 current & last
manufacturer_key	hierarchial warehouse key	NA	NA	NA
manufacturer_code	analytical business attribute	seldom	all history	Type 2
manufacturer_desc	analytical business attribute	seldom	all history	Type 2
design_class_key	hierarchial warehouse key	NA	NA	NA
design_class_code	analytical business attribute	seldom	all history	Type 2
design_class_desc	analytical business attribute	seldom	all history	Type 2



Outline for This Session

- Transformation Process
- Dimension Table Processing
- Dimension Update Types
- **Fact Table Processing**
- Loading Process
- Case Study



Fact Table and Transaction Dimension Processing

- Volume makes **efficiency, performance, and scalability** even more important than with dimension tables
- **Relatively straight forward** process except for sophisticated calculations
- Several methods
- Pick the **best one for your situation**

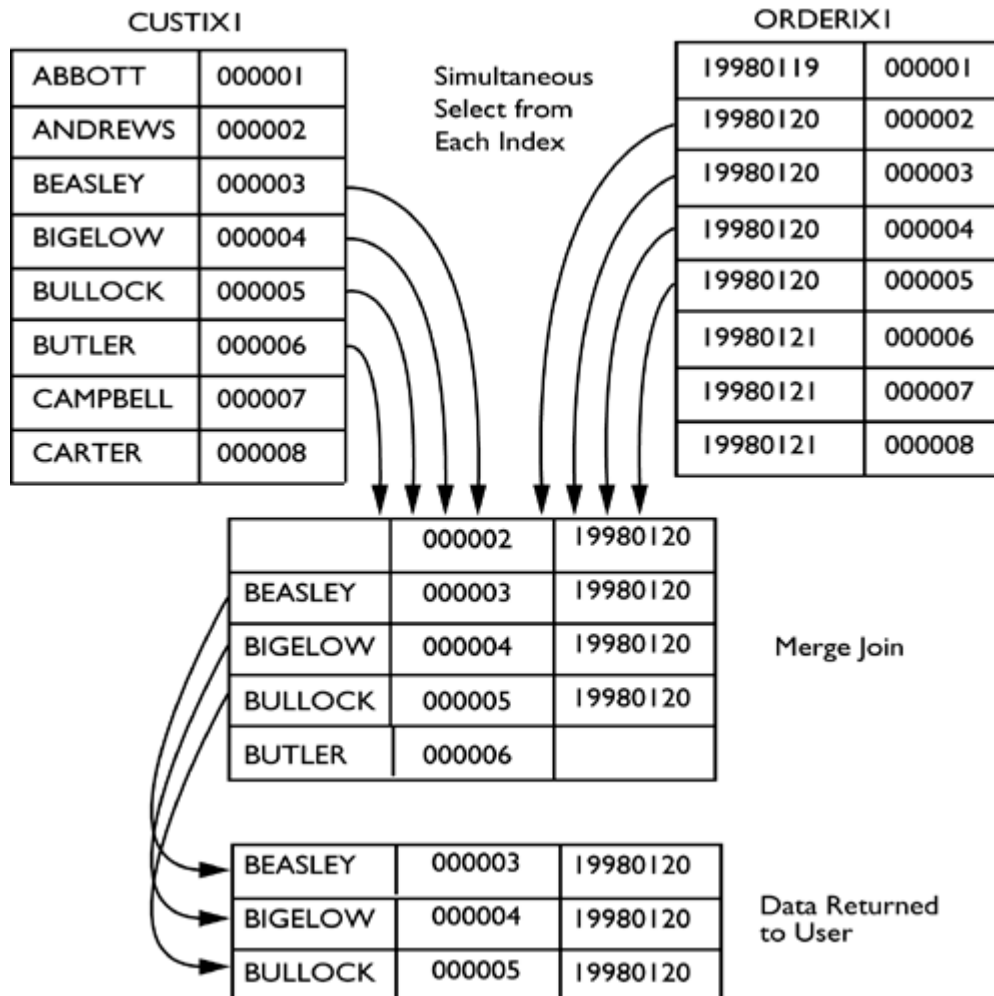


Dimension Key Lookup Methods

- Use **operational keys** to look up the dimension key for the each respective dimension
- **Sort merge**
 - Create a lookup table or file with the operational key and the assigned warehouse key sorted in operational key order
- **No No No**
 - For each fact table row select dimkey from dim where $d.optkey = f.optkey$ (**Cursor Based**)



Dimension Key Lookup Methods



<http://etutorials.org/Misc/advanced+dba+certification+guide+and+reference/Chapter+6.+The+DB2+Optimizer/Joining+in+DB2+UDB/>

Divide and Conquer

- For extremely large volumes
- Parallelize by breaking up input into multiple files
- Parallelize by assigning serial key and break into multiple paths
 - Key + measures
 - Key + dimkey1
 - -----
 - Key + dimkeyn
 - Merge results



Mix Methods

- Parallelize input by **multiple files**
- Sort by **largest dimension**
- Do **hash lookups** for smaller dimensions

Lookup Failures in Fact Table Loads

- Choice 1: Reject fact table rows containing **invalid product ids** until they have been added to the dimension table
- Choice 2: Add row to fact table with default product key signifying **invalid product id**
 - Write error row
- Choice 3: **Add new dimension record** with new key for new product id with all other attributes set to unknown (use with caution)



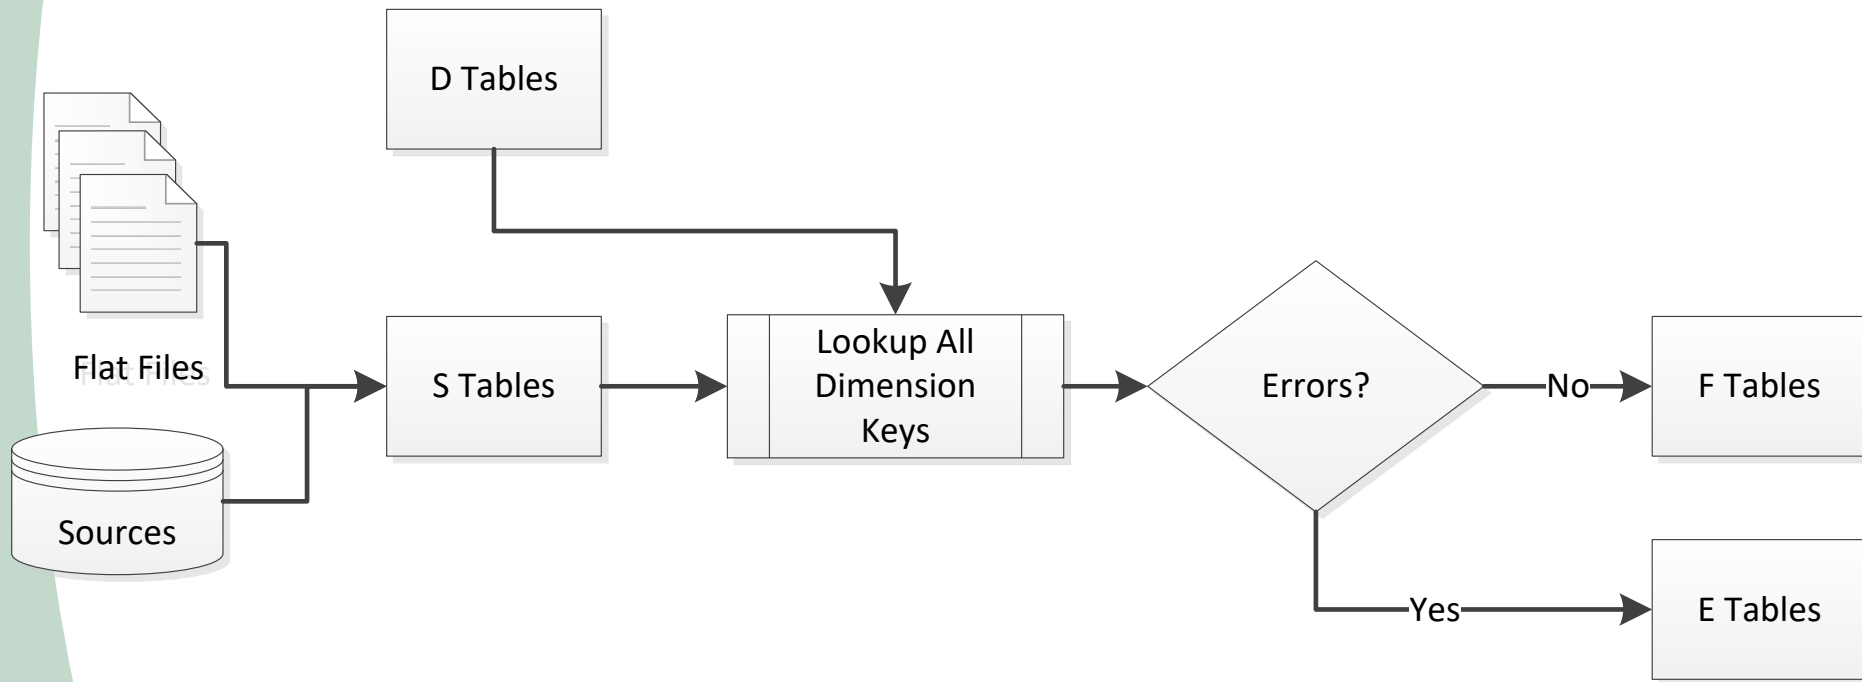
Processing the Error Rows

- Fix the source problem
 - Usually valid code not added yet
 - Or erroneous code entered at data entry
- Add a **new row to the dimension** after the correct code and data are available
- If you used a **key for an invalid row**, do you go back and update the fact table after the dimension is updated?

Lookup Failure Consequences

- Incomplete data until resolved
- Requires fact table update to correct
- Requires dimension update when corrected (can be very dangerous)

Source to Fact Table Processing



Publish Data Quality

- Total Sales figures for 9/5/2009 contain:
 - 2% invalid product codes
 - 0.5% invalid salesperson codes
 - 5% invalid geography codes
- Northwest regional sales for 9/5/2009 contain:
0.5% invalid product codes
- 2% of Product Sales figures for 9/5/2009 are due to invalid products



Outline for This Session

- Transformation Process
- Dimension Table Processing
- Dimension Update Types
- Fact Table Processing
- **Loading Process**
- Case Study



Data Loading

- Data loading takes the prepared data, applies it to the data warehouse, and stores it in the database
- Terminology:
 - Initial Load — populating all the data warehouse tables for the very first time
 - Incremental Load — applying ongoing changes as necessary in a periodic manner
 - Full Refresh — completely erasing the contents of one or more tables and reloading with fresh data (initial load is a refresh of all the tables)

Applying Data: Techniques and Processes

- Load
- Append
- Destructive Merge
- Constructive Merge



Load

- If the target table to be loaded **already exists** and data exists in the table, the load process **wipes out the existing data** and applies the data from the incoming file.
- If the table **is already empty** before loading, the load process simply **applies the data** from the incoming file.

BEFORE

AFTER

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Load

WAREHOUSE	
Key	Data
555	PPPPP
666	QQQQ
777	HHHH

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Append

WAREHOUSE	
Key	Data
111	PPPPP

WAREHOUSE	
Key	Data
111	PPPPP
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Destructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Constructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA*
123	PPPPP
234	BBBBB
345	CCCCC

Append

- **Extension** of the load.
- If data already exists in the table, the append process unconditionally **adds the incoming data**, preserving the existing data in the target table.
- When an incoming record is a **duplicate** of an already existing record, you may define how to handle an incoming duplicate:
 - The incoming record may be allowed to be **added as a duplicate**.
 - In the other option, the incoming duplicate record may be **rejected** during the append process.



BEFORE

AFTER

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Load

WAREHOUSE	
Key	Data
555	PPPPP
666	QQQQ
777	HHHH

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Append

WAREHOUSE	
Key	Data
111	PPPPP

WAREHOUSE	
Key	Data
111	PPPPP
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Destructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Constructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA*
123	PPPPP
234	BBBBB
345	CCCCC

Destructive Merge

- Applies incoming data to the target data.
- If the primary key of an incoming record matches with the key of an existing record, **update the matching** target record.
- If the incoming record is a new record without a match with any existing record, **add** the incoming record to the target table.

BEFORE

AFTER

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Load

WAREHOUSE	
Key	Data
555	PPPPP
666	QQQQ
777	HHHH

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Append

WAREHOUSE	
Key	Data
111	PPPPP

WAREHOUSE	
Key	Data
111	PPPPP
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Destructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Constructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA*
123	PPPPP
234	BBBBB
345	CCCCC

Constructive Merge

- Slightly different from the destructive merge.
- If the primary key of an incoming record matches with the key of an existing record, **leave** the existing record, **add** the incoming record, and mark the added record as **superseding** the old record.

BEFORE

AFTER

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Load

WAREHOUSE	
Key	Data
555	PPPPP
666	QQQQ
777	HHHH

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Append

WAREHOUSE	
Key	Data
111	PPPPP

WAREHOUSE	
Key	Data
111	PPPPP
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Destructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

DATA STAGING	
Key	Data
123	AAAAA
234	BBBBB
345	CCCCC

Constructive Merge

WAREHOUSE	
Key	Data
123	PPPPP

WAREHOUSE	
Key	Data
123	AAAAA*
123	PPPPP
234	BBBBB
345	CCCCC

Outline for This Session

- Transformation Process
- Dimension Table Processing
- Dimension Update Types
- Fact Table Processing
- Loading Process
- **Case Study**



Case Study 4

- ETL Case Walkthrough using R





NATIONAL ENGINEERING CENTER

University of the Philippines
Diliman, Quezon City



Overview of Training Modules of UP NEC

<https://www.facebook.com/upnecanalytics>

<http://www.upnec.com/>

6 Modules

- Analyst Level
 1. Introduction to Business Intelligence
 2. Data Warehousing
 3. Data Mining
- Professional Level
 4. Optimization
 5. Forecasting and Time Series Analysis
 6. R for BI



Analyst Level Certifications

	Data Mining Analyst	Data Warehousing Analyst	Business Intelligence Analyst
Introduction to BI	√	√	√
Data Warehousing		√	√
Data Mining	√		√
Exam	√	√	√



Certification Exam

- Coverage:
 - BI Analyst: All Three Modules
 - Data Mining Analyst: Module 1 and Module 3
 - Data Warehousing Analyst: Module 1 and Module 2
- Type
 - Multiple Choice, Concept Based Questions (From Notes)
 - 3 Hours for BI Analyst
 - 2 Hours for Data Mining and Data Warehousing Analyst
- Offered Monthly c/o Ms. Rea



Professional Level Certifications

	Data Mining Professional	Business Intelligence Professional
Introduction to BI	√	√
Data Warehousing		√
Data Mining	√	√
Optimization		√
Time Series and Forecasting	√	√
Advanced Data Mining	√	√
Exam	√	√



References

- Kimball, Ralph, Margy Ross, Warren Thornthwaite, Joy Mundy, and Bob Becker, *The Data Warehouse Life Cycle Toolkit, Second Edition*, Wiley, 2008, ISBN 978-0-470-14977-5
- Schmitz, Michael D. UCI Irvine Data Warehousing Notes (2014), High Performance Data Warehousing
- Simon, Alan. CIS 391 PPT Slides
- Jeltema ,Bernie, UCI Irvine Data Warehousing Notes (2014), Strategic Frameworks, Inc.

